

Nirdizati: A Web-Based Tool for Predictive Process Monitoring

Kerwin Jorbina¹, Andrii Rozumnyi¹, Ilya Verenich^{1,2}, Chiara Di Francescomarino³, Marlon Dumas¹, Chiara Ghidini³, Fabrizio Maria Maggi¹, Marcello La Rosa³, and Simon Raboczi³

¹ University of Tartu, Estonia

{kerwin.jorbina, andriiro, marlon.dumas, f.m.maggi}@ut.ee

² Queensland University of Technology, Australia

{ilya.verenich, m.larosa, simon.raboczi}@qut.edu.au

³ FBK IRST, Trento, Italy

{dfmchiara, ghidini}@fbk.eu

Abstract. This paper introduces Nirdizati: A web-based application for generating predictions about running cases of a business process. Nirdizati is a configurable full-stack web application that supports users in selecting and tuning prediction methods from a list of implemented algorithms and enables the continuous prediction of various performance indicators at runtime. The tool can be used to predict the outcome, the next events, or the remaining time of each case of a process. For example, in a lead-to-order process, Nirdizati can predict which customer leads will convert to purchase orders and when. In a claims handling process, it can predict if a claim decision will be made on time or late. The results of the predictions, as well as the real-time summary statistics about the process executions, are presented in a dashboard that offers multiple visualization options. Based on these predictions, process participants can act proactively to resolve or mitigate potential process performance violations. The target audience of this demonstration includes process mining researchers as well as practitioners interested in exploring the potential of predictive process monitoring.

Keywords: Predictive Process Monitoring, Machine Learning

1 Introduction

Predictive Process Monitoring [1] is an emerging paradigm based on the continuous generation of predictions about the future values of user-specified performance indicators of a currently running process execution. In this paradigm, a user defines the type of predictions she is interested in and provides a set of historical execution traces. Based on the analysis of these traces, the idea of predictive monitoring is to continuously provide the user with predictions and estimated values of the measure of interest. Such predictions generally depend both on: (i) the sequence of activities executed in a given case; and (ii) the values of data attributes after each activity execution in the case.

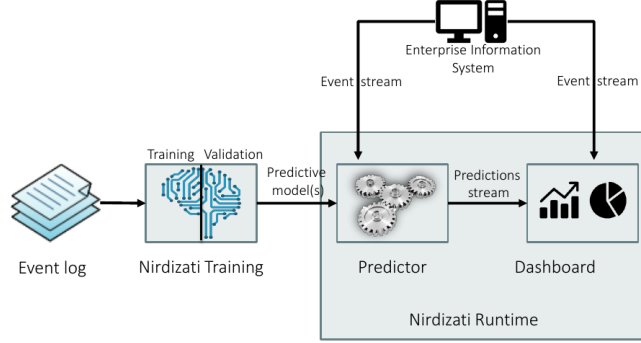


Fig. 1. High-level overview of Nirdizati.

As an example, consider a business process for managing loan applications; the applicant can be advised about the combinations of loan amount and duration that are the most likely to lead to an acceptance of the application, given some contextual information about the application and the personal data of the applicant (e.g., age, salary, etc.).

Several approaches have been proposed in the literature to tackle common predictive process monitoring tasks. However, so far, these approaches have largely remained in the academic domain and have not been widely applied in real-time scenarios, in which users require a continuous predictive support.

In this paper, we present Nirdizati, a pioneering open-source Web-based predictive monitoring tool, which is able to fill this gap between research and practice, by providing business analysts with a highly flexible instrument for the selection, generation and analysis of different predictive models, and end-users with continuous runtime predictions.

Nirdizati consists of two components: *Nirdizati Training* and *Nirdizati Runtime* (Fig. 1). Nirdizati Training takes as input a business process event log and produces one or more prediction models, which can then be deployed in Nirdizati Runtime. Once a model is deployed, Nirdizati Runtime listens to a stream of events related to the process, and produces a stream of predictions. These predictions are then visualized in a continuously updated Web dashboard.

2 Nirdizati Training

Nirdizati Training is the component of Nirdizati that allows users to produce predictive models later used by Nirdizati Runtime for making predictions on streams of events. It provides several algorithms for generating predictive models suitable for different types of predictions and tailored to each specific dataset. For example, it is able to build predictive models for predicting remaining time, the next activity to be performed, whether a certain outcome will be achieved or not and even the overall workload per day. To this aim, the training component of




Fig. 2. Front-end of Nirdizati Training.

Nirdizati relies on two phases: a training and a validation phase. In the former, one or more predictive models are built; in the latter, their suitability to the specific dataset is evaluated, so as to support the user in selecting the predictive model that ensures the best results.

Nirdizati Training is composed of a front-end application (Fig. 2), which allows users to configure the predictive models and to assess the goodness-of-fit of the built models, and a back-end application responsible for the actual training and validation. The back-end application is in turn composed of four submodules shown in Fig. 3. A Log Manager is in charge of managing the logs. Uploading and retrieving the logs are the basic operations of this module. The Encoder is responsible for parsing the logs, labeling them according to the desired type of prediction, and preparing the data for the training phase. In the Training submodule, the encoded data is split into training and validation set used for

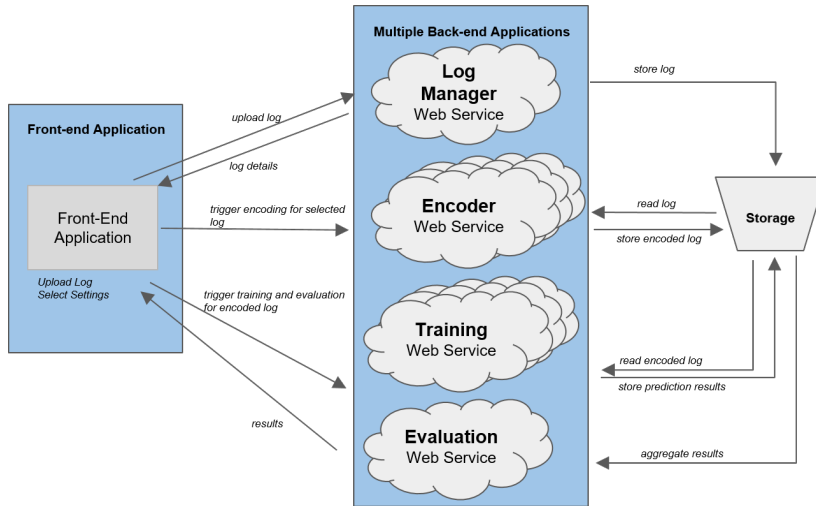


Fig. 3. High-level overview of Nirdizati Training.




Fig. 4. High-level data flow diagram of Nirdizati Runtime.

evaluation purposes, and the predictive models are built from the training data. Finally, the Evaluation submodule tests the validation data against the model(s) created to get accuracy measures with respect to the ground truth, available from the complete traces in the validation data. The back-end application comes with a storage module for saving the uploaded logs and the built predictive models.

3 Nirdizati Runtime

Once the predictive models have been trained, they are used by the Runtime component to make predictions on ongoing cases. As input, Nirdizati Runtime takes a stream of events produced by information systems that register the tasks performed in an organization (Fig. 4). The stream is processed by a web server. In particular, the web server communicates with a predictive engine that queries the predictive models obtained by Nirdizati Training to get predictions about all the cases recorded in the event stream. After the web server receives the predictions from each predictive model, it stores them into a database. From the user’s side, the results are visualized in a dashboard-based interface, which supports several visualization options.

When a user logs into the system, she initially sees a table listing both currently ongoing cases (colored in gray) as well as completed cases (colored in green), as shown in Fig. 5. For each case, Nirdizati Runtime provides a range of summary statistics including the number of events in a given case, starting time of a case and the time when the latest event has occurred. For the ongoing cases, Nirdizati Runtime also estimates their completion time and a range of other process-specific indicators, such as the probability that a customer will not accept an offer from a salesperson. For completed cases, instead, it shows the actual completion time and the actual values of the indicators.

Process workers and operational managers – typical users of Nirdizati – can set the process performance targets and subscribe to a stream of warnings and alerts generated whenever these targets are predicted to be violated. Thus, they

NIRDIZATI

Loan requests

RUNNING CASES

208

COMPLETED CASES

3

COMPLETED EVENTS

1,568

AVG CASE LENGTH

9.33

AVG CASE DURATION

21h 23m 40s

Detail view

Outcomes

Case duration

Remaining time

Case length

Search

case ID	Completed	Events elapsed	Start time	Latest event time	Predicted/Actual completion time	Predicted/Actual duration	Probability to be slow	Probability to be rejected
c_2026841022	false	9	Oct 05 2016 09:07:17	Oct 06 2016 08:05:44	Oct 29 2016 11:49:42	24d 2h 42m	0.57	0.804
c_1243522820	false	8	Oct 01 2016 10:12:27	Oct 01 2016 10:19:56	Oct 22 2016 07:26:47	20d 21h 14m	0.354	0.201
c_1811993938	true	9	Oct 01 2016 10:22:34	Oct 01 2016 10:46:14	Oct 01 2016 10:46:14	0h 23m 40s	-	-
c_1080763304	false	9	Oct 01 2016 10:59:16	Oct 01 2016 13:59:36	Oct 26 2016 00:09:03	24d 13h 9m	0.597	0.277
c_1460054354	false	11	Oct 01 2016 11:36:27	Oct 01 2016 11:46:11	Oct 27 2016 02:46:51	25d 15h 10m	0.662	0.227

Showing 1 to 5 of 210 rows

5 rows per page

<

1

2

3

4

5

...

42

>

Fig. 5. Main view of Nirdizati Runtime.

will be capable of making informed, data-driven decisions to get a better control of the process executions. This is especially beneficial for processes where process participants have more leeway to make corrective actions (for example, in a lead management process).

4 Conclusion

Nirdizati is a configurable full-stack web application that supports users in selecting the preferred prediction method from a list of implemented methods and enables the continuous prediction of various performance indicators at runtime. The results of the predictions, as well as the real-time summary statistics about the process executions, are presented in a dashboard that offers multiple visualization options.

A video demo of Nirdizati can be found at <http://youtu.be/0nr14lX04-I>. A public release of Nirdizati Training and Nirdizati Runtime are available at <http://training.nirdizati.com> and <http://dashboard.nirdizati.com>, respectively. The source code is available under the Lesser GNU Public License (LGPL) at <http://github.com/nirdizati>.

References

1. Maggi, F., Di Francescomarino, C., Dumas, M., Ghidini, C.: Predictive monitoring of business processes. In: Proc. of the International Conference on Advanced Information Systems Engineering (CAiSE). Springer (2013)